

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
from sklearn import metrics
```

In [2]:

```
df=pd.read_csv('after_week1.csv')
df.head()
```

Out[2]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFun
0	6	148.0	72.0	35.000000	79.799479	33.6	
1	1	85.0	66.0	29.000000	79.799479	26.6	
2	8	183.0	64.0	20.536458	79.799479	23.3	
3	1	89.0	66.0	23.000000	94.000000	28.1	
4	0	137.0	40.0	35.000000	168.000000	43.1	

In [3]:

```
x=df.iloc[:, :-1].values
y=df.iloc[:, -1].values
```

In [4]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=0)
```

In [5]:

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(614L, 8L)
(154L, 8L)
(614L,)
(154L,)
```

In [6]:

```
from sklearn.preprocessing import StandardScaler
```

In [7]:

```
Scale=StandardScaler()  
x_train_std=Scale.fit_transform(x_train)  
x_test_std=Scale.transform(x_test)
```

In [8]:

```
norm=lambda a:(a-min(a))/(max(a)-min(a))
```

In [9]:

```
df_norm=df.iloc[:, :-1]
```

In [10]:

```
df_normalized=df_norm.apply(norm)
```

In [11]:

```
x_train_norm,x_test_norm,y_train_norm,y_test_norm=train_test_split(df_normalized.values,y,t
```

In [12]:

```
print(x_train_norm.shape)  
print(x_test_norm.shape)  
print(y_train_norm.shape)  
print(y_test_norm.shape)
```

```
(614L, 8L)  
(154L, 8L)  
(614L,)  
(154L,)
```

In [13]:

```
from sklearn.neighbors import KNeighborsClassifier  
knn_model = KNeighborsClassifier(n_neighbors=25)  
#Using 25 Neighbors just as thumb rule sqrt of observation  
knn_model.fit(x_train_std,y_train)  
knn_pred=knn_model.predict(x_test_std)
```

In [14]:

```

print("Model Validation ==>\n")
print("Accuracy Score of KNN Model::")
print(metrics.accuracy_score(y_test,knn_pred))
print("\n","Classification Report::")
print(metrics.classification_report(y_test,knn_pred),'\n')
print("\n","ROC Curve")
knn_prob=knn_model.predict_proba(x_test_std)
knn_prob1=knn_prob[:,1]
fpr,tpr,thresh=metrics.roc_curve(y_test,knn_prob1)
roc_auc_knn=metrics.auc(fpr,tpr)
plt.figure(dpi=80)
plt.title("ROC Curve")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.plot(fpr,tpr,'b',label='AUC Score = %0.2f'%roc_auc_knn)
plt.plot(fpr,fpr,'r--',color='red')
plt.legend()

```

Model Validation ==>

Accuracy Score of KNN Model::

0.8181818181818182


('\n', 'Classification Report::')

(u'	precision	recall	f1-score	support	0
0.85	0.90	0.87	107	1	0.73
0.68	47	micro avg	0.82	0.82	0.82
macro avg	0.79	0.77	0.78	154	nweighted avg
0.82	0.81	154	n'	n'	0.81

('\n', 'ROC Curve')

Out[14]:

<matplotlib.legend.Legend at 0x12b649e8>



In [15]:

```
from sklearn.neighbors import KNeighborsClassifier
knn_model_norm = KNeighborsClassifier(n_neighbors=25)
#Using 25 Neighbors just as thumb rule sqrt of observation
knn_model_norm.fit(x_train_norm,y_train_norm)
knn_pred_norm=knn_model_norm.predict(x_test_norm)
```

In [16]:

```

print("Model Validation ==>\n")
print("Accuracy Score of KNN Model with Normalization::")
print(metrics.accuracy_score(y_test_norm,knn_pred_norm))
print("\n","Classification Report::")
print(metrics.classification_report(y_test_norm,knn_pred_norm),'\n')
print("\n","ROC Curve")
knn_prob_norm=knn_model.predict_proba(x_test_norm)
knn_prob_norm1=knn_prob_norm[:,1]
fpr,tpr,thresh=metrics.roc_curve(y_test_norm,knn_prob_norm1)
roc_auc_knn=metrics.auc(fpr,tpr)
plt.figure(dpi=80)
plt.title("ROC Curve")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.plot(fpr,tpr,'b',label='AUC Score = %0.2f'%roc_auc_knn)
plt.plot(fpr,fpr,'r--',color='red')
plt.legend()

```

Model Validation ==>

Accuracy Score of KNN Model with Normalization::

0.8311688311688312


('\n', 'Classification Report::')

	precision	recall	f1-score	support
0	0.86	0.90	0.88	107
1	0.71	0.74	0.72	47
micro avg	0.83	0.83	0.83	154
macro avg	0.80	0.79	0.80	154
weighted avg	0.83	0.83	0.83	154

('\n', 'ROC Curve')

Out[16]:

<matplotlib.legend.Legend at 0x12c82e10>



In [17]:

```

from sklearn.svm import SVC
svc_model_linear = SVC(kernel='linear',random_state=0,probability=True,C=0.01)
svc_model_linear.fit(x_train_std,y_train)
svc_pred=svc_model_linear.predict(x_test_std)

```

In [18]:

```

print("Model Validation ==>\n")
print("Accuracy Score of SVC Model with Linear Kernel::")
print(metrics.accuracy_score(y_test,svc_pred))
print("\n","Classification Report::")
print(metrics.classification_report(y_test,svc_pred),'\n')
print("\n","ROC Curve")
svc_prob_linear=svc_model_linear.predict_proba(x_test_std)
svc_prob_linear1=svc_prob_linear[:,1]
fpr,tpr,thresh=metrics.roc_curve(y_test,svc_prob_linear1)
roc_auc_svc=metrics.auc(fpr,tpr)
plt.figure(dpi=80)
plt.title("ROC Curve")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.plot(fpr,tpr,'b',label='AUC Score = %0.2f'%roc_auc_svc)
plt.plot(fpr,fpr,'r--',color='red')
plt.legend()

```

Model Validation ==>

Accuracy Score of SVC Model with Linear Kernel::

0.8116883116883117


('\n', 'Classification Report::')

	precision	recall	f1-score	support
0.83	0.92	0.87	107	1
0.65	0.47	0.79	0.81	154
macro avg	0.79	0.75	0.76	154
0.81	0.80	0.81	0.81	154

('\n', 'ROC Curve')

Out[18]:

<matplotlib.legend.Legend at 0x130cd0b8>



In [19]:

```

from sklearn.svm import SVC
svc_model_rbf = SVC(kernel='rbf',random_state=0,probability=True,C=1)
svc_model_rbf.fit(x_train_std,y_train)
svc_pred_rbf=svc_model_rbf.predict(x_test_std)

```

In [20]:

```

print("Model Validation ==>\n")
print("Accuracy Score of SVC Model with RBF Kernel::")
print(metrics.accuracy_score(y_test,svc_pred_rbf))
print("\n","Classification Report::")
print(metrics.classification_report(y_test,svc_pred_rbf),'\n')
print("\n","ROC Curve")
svc_prob_rbf=svc_model_linear.predict_proba(x_test_std)
svc_prob_rbf1=svc_prob_rbf[:,1]
fpr,tpr,thresh=metrics.roc_curve(y_test,svc_prob_rbf1)
roc_auc_svc=metrics.auc(fpr,tpr)
plt.figure(dpi=80)
plt.title("ROC Curve")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.plot(fpr,tpr,'b',label='AUC Score = %0.2f'%roc_auc_svc)
plt.plot(fpr,fpr,'r--',color='red')
plt.legend()

```

Model Validation ==>

Accuracy Score of SVC Model with RBF Kernel::

0.7727272727272727


('\n', 'Classification Report::')

	precision	recall	f1-score	support
0	0.81	0.88	0.84	107
1	0.59	0.47	0.53	47
micro avg	0.77	0.77	0.77	154
macro avg	0.73	0.71	0.72	154
weighted avg	0.77	0.77	0.76	154

('\n', 'ROC Curve')

Out[20]:

<matplotlib.legend.Legend at 0x131f73c8>



In [21]:

```

from sklearn.linear_model import LogisticRegression
lr_model = LogisticRegression(C=0.01)
lr_model.fit(x_train_std,y_train)
lr_pred=lr_model.predict(x_test_std)

```

In [22]:

```

print("Model Validation ==>\n")
print("Accuracy Score of Logistic Regression Model::")
print(metrics.accuracy_score(y_test,lr_pred))
print("\n","Classification Report::")
print(metrics.classification_report(y_test,lr_pred),'\n')
print("\n","ROC Curve")
lr_prob=lr_model.predict_proba(x_test_std)
lr_prob1=lr_prob[:,1]
fpr, tpr, thresh=metrics.roc_curve(y_test,lr_prob1)
roc_auc_lr=metrics.auc(fpr,tpr)
plt.figure(dpi=80)
plt.title("ROC Curve")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.plot(fpr,tpr,'b',label='AUC Score = %0.2f'%roc_auc_lr)
plt.plot(fpr,fpr,'r--',color='red')
plt.legend()

```

Model Validation ==>

Accuracy Score of Logistic Regression Model::

0.7922077922077922


('\n', 'Classification Report::')

	precision	recall	f1-score	support
0.84	0.87	0.85	107	1
0.64	0.47	0.62	154	154
micro avg	0.76	0.74	0.75	154
macro avg	0.79	0.79	0.79	154
weighted avg	0.79	0.79	0.79	154

('\n', 'ROC Curve')

Out[22]:

<matplotlib.legend.Legend at 0x13498860>



In [23]:

```
from sklearn.ensemble import RandomForestClassifier
rf_model = RandomForestClassifier(n_estimators=1000,random_state=0)
rf_model.fit(x_train_std,y_train)
rf_pred=rf_model.predict(x_test_std)
```

In [24]:

```
print("Model Validation ==>\n")
print("Accuracy Score of Logistic Regression Model::")
print(metrics.accuracy_score(y_test,rf_pred))
print("\n","Classification Report::")
print(metrics.classification_report(y_test,rf_pred),'\n')
print("\n","ROC Curve")
rf_prob=rf_model.predict_proba(x_test_std)
rf_prob1=rf_prob[:,1]
fpr,tpr,thresh=metrics.roc_curve(y_test,rf_prob1)
roc_auc_rf=metrics.auc(fpr,tpr)
plt.figure(dpi=80)
plt.plot(fpr,tpr,'b',label='AUC Score = %0.2f'%roc_auc_rf)
plt.title("ROC Curve")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.plot(fpr,fpr,'r--',color='red')
plt.legend()
```

Model Validation ==>

Accuracy Score of Logistic Regression Model::

0.8246753246753247


('\n', 'Classification Report::')

	precision	recall	f1-score	support	0	1	weighted avg
0.88	0.87	0.87	0.87	107	0.71	0.72	
0.72	0.47	0.72	0.57	154	0.82	0.82	0.8
micro avg	0.79	0.80	0.79	154			
macro avg	0.67	0.79	0.73	154			
weighted avg	0.82	0.83	0.83	154			

('\n', 'ROC Curve')

Out[24]:

<matplotlib.legend.Legend at 0x16957fd0>



In []: